

# NSC Final Lab - NAT Implementation using P4 language with bmv2 switch

---

## Repository link:

All the code and environment setup can be found in the repository : [https://github.com/Zhih25/NSC\\_final](https://github.com/Zhih25/NSC_final)

## Introduction:

This project is a simple implementation of Layer 4 port based NAT (Network Address Translation) using P4 language with bmv2 switch. The environment will simulate a scheme with internal network and outer network, then any packet with L4 TCP/UDP header will be translated to the outer/internal network with a public IP address using the src/dst transport layer port.

## Motivation:

Because of my CS project is related to the P4 language, so actually I have learned P4 for a while. During the time learning P4, I found that I am interested in the P4 language, also I found that I was not much familiar with the concept of NAT, so I choose to implement the NAT in P4 language to practice the basic concept of P4 language.

## Main Concept:

Any packet from the internal network to the outer network will be assign an unique port number using the combination of source IP and source port to avoid the conflict. The packet will be translated to the public IP address, and the reply packet from the outer network will use that unique port number to trace back the corresponding internal IP and port number.


## Expected Result:

1. Switch forward the packet between the internal network and the outer network correctly.
2. Switch is able to handle TCP and UDP packet.
3. Switch is able to assign a unique port number used in public IP address for a internal packet and change the source IP address to the public IP address.
4. Switch is able to translate the unique port number back to the corresponding internal IP and port number.
5. Switch can be statically set up to support the sever running in the internal network.

## Network Topology

The topology is a network with 4 hosts and 1 switch, I simply divide the network into three parts, internal network h1 and h2, two outer network h3 and h4, three networks are connected by the switch s1. The MAC

address and IP address are shown in the following table:



Network 1:

- h1 and h2, with MAC address prefix **08:00:00:01:{host\_number}** and IP prefix **10.0.1.{host\_number}/24**
- Default gateway: **10.0.1.10** at eth0 to switch s1
- Using public IP address: **140.113.0.1**

Network 2:

- h3 with MAC address **08:00:00:03:03** and IP address **140.113.100.100/24**
- Default gateway: **140.113.100.50** at eth0 to switch s1

Network 3:

- h4 with MAC address **08:00:00:04:04** and IP address **140.113.200.200/24**
- Default gateway: **140.113.200.50** at eth0 to switch s1

Host	MAC Address	IP Address
h1	00:00:00:00:01:01	10.0.1.1/24
h2	00:00:00:00:01:02	10.0.1.2/24

Host	MAC Address	IP Address
h3	00:00:00:00:03:03	140.113.100.100/24
h4	00:00:00:00:04:04	140.113.200.200/24

## Experiment Result:

I design 5 tests to test the if functionality of the NAT switch can fulfill the expected result. The result is shown below, and the testing flow has been shown in the [README](#) file.

### 1. Forwarding between the internal network.

This part is aimed to test the connection between the internal network, the packet should be forwarded correctly between h1 and h2.

Result:

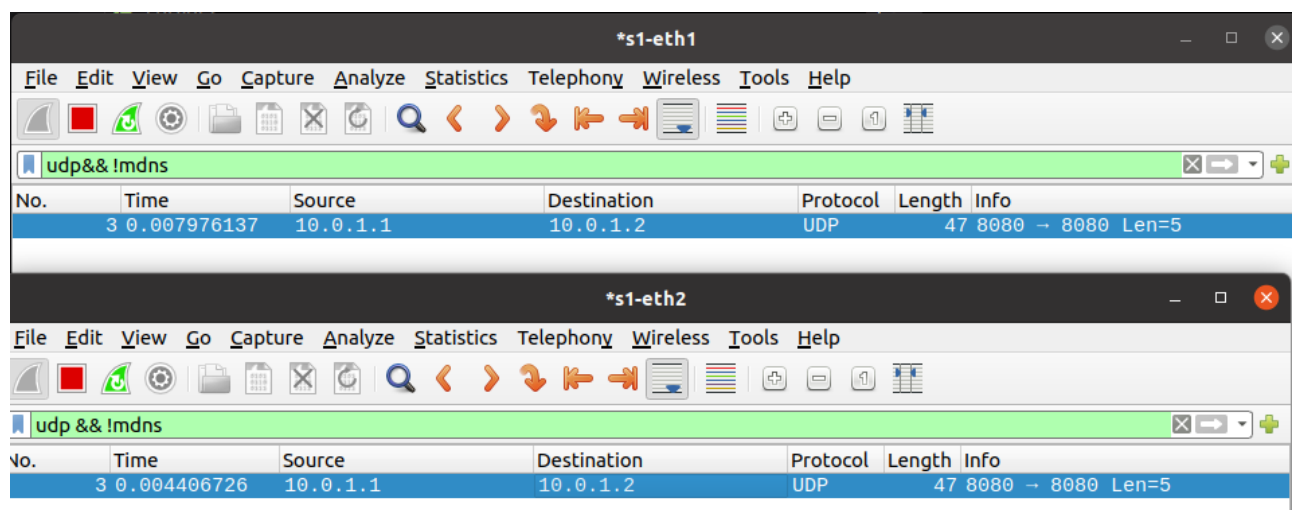
#### h1's terminal

```
"Node: h1"
root@zhi-p4:/home/zhi/Desktop/nsc_final/src/nat# python3 udp.py 1
UDP socket bound to 10.0.1.1:8080
Enter destination IP address: 10.0.1.2
Enter destination port: 8080
Enter message to send: hello
Enter destination IP address: "
```

#### h2's terminal

```
"Node: h2"
root@zhi-p4:/home/zhi/Desktop/nsc_final/src/nat# python3 udp.py 2
UDP socket bound to 10.0.1.2:8080
Enter destination IP address:
Received message from ('10.0.1.1', 8080): hello
"
```

Wireshark in h1 and h2:



In the screenshot above, we can see that the packet is forwarded correctly from h1 to h2.

### 2. Forwarding between the external network.

This part is aimed to test the connection between the external network, the packet should be forwarded correctly between h3 and h4, both IP and port are not rewrite by the switch.

Result:

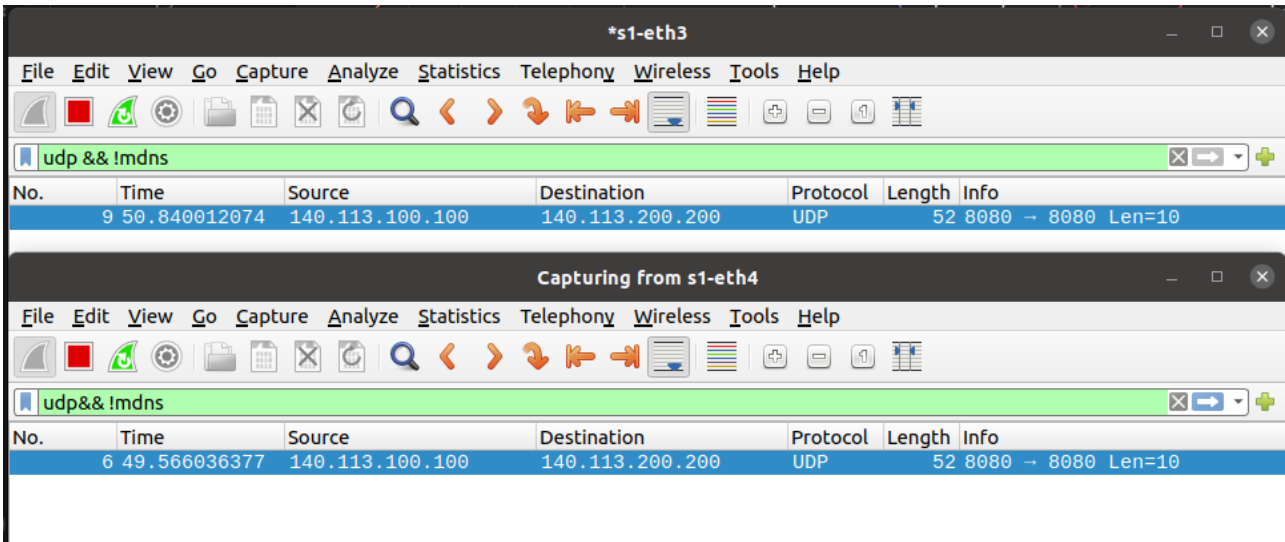
**h3's terminal**

```
"Node: h3"
root@zhi-p4:/home/zhi/Desktop/nsc_final/src/nat# python3 udp.py 3
UDP socket bound to 140.113.100.100:8080
Enter destination IP address: 140.113.200.200
Enter destination port: 8080
Enter message to send: hi from h3
Enter destination IP address: 
```

**h4's terminal**

```
"Node: h4"
root@zhi-p4:/home/zhi/Desktop/nsc_final/src/nat# python3 udp.py 4
UDP socket bound to 140.113.200.200:8080
Enter destination IP address:
Received message from ('140.113.100.100', 8080): hi from h3
```

Wireshark in h3 and h4:



In the screenshot above, we can see that the packet is forwarded correctly from h3 to h4, both IP and port are not rewrite by the switch.

3. Test the UDP connection from an internal network host to an outer network host.

This part is going to test the UDP connection from h2 to h3. Because the packet is sent from the internal network to the outer network, NAT should be applied, the source IP address and port number of the packet should be translated to the public IP address 140.113.0.1 and a unique port number assigned by the switch.

Result:

**h2's terminal**

**h3's terminal**

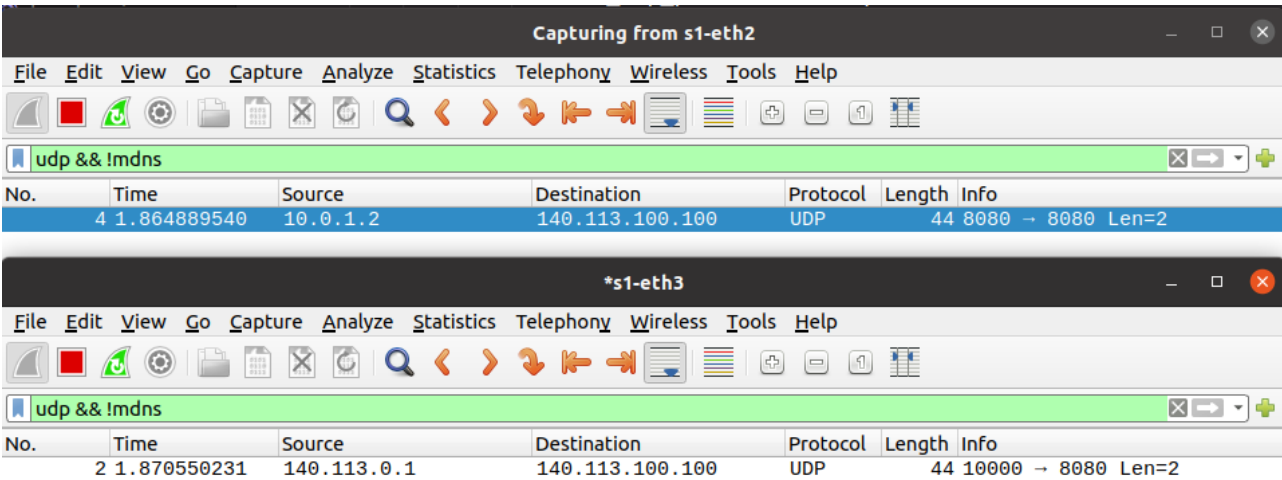
h2's terminal

```
"Node: h2"
root@zhi-p4:/home/zhi/Desktop/nsc_final/src/nat# python3 udp.py 2
UDP socket bound to 10.0.1.2:8080
Enter destination IP address: 140.113.100.100
Enter destination port: 8080
Enter message to send: hi
Enter destination IP address: []
```

h3's terminal

```
"Node: h3"
root@zhi-p4:/home/zhi/Desktop/nsc_final/src/nat# python3 udp.py 3
UDP socket bound to 140.113.100.100:8080
Enter destination IP address:
Received message from ('140.113.0.1', 10000): hi
[]
```

Wireshark in h2 and h3:



From the screenshot above, we can see that the IP and source port shown in h2 is 10.0.1.2 and 8080, but when the packet is received by h3 is the public IP address 140.113.0.1 and the unique port number 10000 after the NAT translation. Thus the packet is forwarded correctly from h2 to h3.

4. TCP connection from two host in same internal network to the outer network host.

This part is going to test if two internal network hosts are trying to connect to a same outer network host, the NAT function will translate the source IP address and port to the public IP address and port.

In the test, both h1 and h2 are binding to the same port 8888 and try to connect to h3, the packet should be translated to the public IP address 140.113.0.1 and the port number should be unique.

Result:

h1's terminal

h2's terminal

h3's terminal

## h1's terminal

```
"Node: h1"
root@h1:~# python3 tcp_sender.py 1
Connected to 140.113.100.100:12345
Enter message to send (or 'ex' to disconnect): hello from h1
Enter message to send (or 'ex' to disconnect): []
```

## h2's terminal

```
"Node: h2"
root@h1:~# python3 tcp_sender.py 2
Connected to 140.113.100.100:12345
Enter message to send (or 'ex' to disconnect): hello from h2
Enter message to send (or 'ex' to disconnect): []
```

## h3's terminal

```
"Node: h3"
root@h1:~# python3 tcp_receiver.py
Waiting for a connection...
Connection from ('140.113.0.1', 10000)
Received message from ('140.113.0.1', 10000): hello from h1
Received message from ('140.113.0.1', 10000): hello from h2
[]
```

Wireshark in h1, h2 and h3: h1's wireshark:

*s1-eth1							
File Edit View Go Capture Analyze Statistics Telephony Wireless Tools Help							
tcp							
No.	Time	Source	Destination	Protocol	Length	Info	
17	44.410794244	10.0.1.1	140.113.100.100	TCP	74	8888 → 12345 [SYN] Seq=0 W	
18	44.419357018	140.113.100.100	10.0.1.1	TCP	74	12345 → 8888 [SYN, ACK] Seq	
19	44.420350974	10.0.1.1	140.113.100.100	TCP	66	8888 → 12345 [ACK] Seq=1 Ac	
20	51.396004558	10.0.1.1	140.113.100.100	TCP	79	8888 → 12345 [PSH, ACK] Seq	
21	51.401712414	140.113.100.100	10.0.1.1	TCP	66	12345 → 8888 [ACK] Seq=1 Ac	

h2's wireshark:

*s1-eth2							
File Edit View Go Capture Analyze Statistics Telephony Wireless Tools Help							
tcp							
No.	Time	Source	Destination	Protocol	Length	Info	
15	37.655218001	10.0.1.2	140.113.100.100	TCP	74	8888 → 12345 [SYN] Seq=0 W	
16	37.660163341	140.113.100.100	10.0.1.2	TCP	74	12345 → 8888 [SYN, ACK] Seq	
17	37.660263470	10.0.1.2	140.113.100.100	TCP	66	8888 → 12345 [ACK] Seq=1 Ac	
19	55.602776209	10.0.1.2	140.113.100.100	TCP	79	8888 → 12345 [PSH, ACK] Seq	
20	55.617164518	140.113.100.100	10.0.1.2	TCP	66	12345 → 8888 [ACK] Seq=1 Ac	

h3's wireshark:

s1-eth3							
File Edit View Go Capture Analyze Statistics Telephony Wireless Tools Help							
tcp							
No.	Time	Source	Destination	Protocol	Length	Info	
4	34.039383066	140.113.0.1	140.113.100.100	TCP	74	10000 → 12345 [SYN] Seq=0 W	
5	34.039403107	140.113.100.100	140.113.0.1	TCP	74	12345 → 10000 [SYN, ACK] Se	
6	34.043747565	140.113.0.1	140.113.100.100	TCP	66	10000 → 12345 [ACK] Seq=1 A	
7	39.515142443	140.113.0.1	140.113.100.100	TCP	74	10001 → 12345 [SYN] Seq=0 W	
8	39.515167894	140.113.100.100	140.113.0.1	TCP	74	12345 → 10001 [SYN, ACK] Se	
9	39.524616124	140.113.0.1	140.113.100.100	TCP	66	10001 → 12345 [ACK] Seq=1 A	
10	46.497182321	140.113.0.1	140.113.100.100	TCP	79	10001 → 12345 [PSH, ACK] Se	
11	46.497201262	140.113.100.100	140.113.0.1	TCP	66	12345 → 10001 [ACK] Seq=1 A	
12	51.988579105	140.113.0.1	140.113.100.100	TCP	79	10000 → 12345 [PSH, ACK] Se	
13	51.988649781	140.113.100.100	140.113.0.1	TCP	66	12345 → 10000 [ACK] Seq=1 A	

From the screenshot above, we can see that TCP connection is working correctly between h1 and h3, and h2 and h3. From the wireshark in h1 and h2, we can see that the source IP address and port number are both their own IP address and port number in the internal network, but when the packet is received by h3, the source IP address is the public IP address **140.113.0.1** and the port number is unique. So the NAT function is working correctly.

5. Test the running server in the internal network can be accessed by the outer network.

In the previous part, once we want to start a connection between the internal network and the outer network, we need to set up the connection from the internal network to the outer network first. But in P4 NAT, the connection can be forwarding from the outer network to the internal network directly, any connection want to access the server via specific port, can be forwarded to the specific host in the internal network.

In the test, the server is running in h1, and the server is listening to the port **80** to simulate a web server, the client is running in h4 and try to access the server in h1. h4 is expected to receive the message from the server in h1.

Result:

h1's terminal

```
"Node: h1"
root@zhi-p4:/home/zhi/Desktop/nsc_final/src/nat# python3 server.py
Waiting for a connection...
Connection from ('140.113.200.200', 32952)
[]
```

h4's terminal

```
"Node: h4"
root@zhi-p4:/home/zhi/Desktop/nsc_final/src/nat# python3 client.py
Connected to 140.113.0.1:80
Hello from server
Hello from server
Hello from server
Hello from server
Hello from server
Hello from server
[]
```

Wireshark in h1:

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000000	140.113.200.200	140.113.0.1	TCP	74	39398 → 80 [SYN] Seq=0 Win=
2	0.007299269	140.113.0.1	140.113.200.200	TCP	74	80 → 39398 [SYN, ACK] Seq=0
3	0.007369794	140.113.200.200	140.113.0.1	TCP	66	39398 → 80 [ACK] Seq=1 Ack=
4	0.024108558	140.113.0.1	140.113.200.200	TCP	83	80 → 39398 [PSH, ACK] Seq=1
5	0.024122059	140.113.200.200	140.113.0.1	TCP	66	39398 → 80 [ACK] Seq=1 Ack=
6	2.015998546	140.113.0.1	140.113.200.200	TCP	83	80 → 39398 [PSH, ACK] Seq=1
7	2.016018338	140.113.200.200	140.113.0.1	TCP	66	39398 → 80 [ACK] Seq=1 Ack=

Code Implementation:

Header

```

header ethernet_t {
    macAddr_t dstAddr;
    macAddr_t srcAddr;
    bit<16> etherType;
}

header ipv4_t {
    bit<4> version;
    bit<4> ihl;
    bit<8> diffserv;
    bit<16> totalLen;
    bit<16> identification;
    bit<3> flags;
    bit<13> fragOffset;
    bit<8> ttl;
    bit<8> protocol;
    bit<16> hdrChecksum;
    ip4Addr_t srcAddr;
    ip4Addr_t dstAddr;
}

header udp_t{
    bit<16> srcPort;
    bit<16> dstPort;
    bit<16> length;
    bit<16> checksum;
}

header tcp_t{
    bit<16> srcPort;
    bit<16> dstPort;
    bit<32> seq_num;
    bit<32> ack_num;
    bit<4> data_offset;
    bit<3> reserved;
    bit<9> ctl_flag;
    bit<16> window_size;
    bit<16> checksum;
    bit<16> urgent_num;
}

header_union l4_t{
    udp_t udp;
    tcp_t tcp;
}

struct metadata {
    bit<16> tcp_length;
    bit<16> tot_length;
    bit<16> udp_length;
}

struct headers {
    ethernet_t ethernet;
    ipv4_t ipv4;
}

```



```

        l4_t      l4;
    }

```

Header struct is set up as the above code, because NAT is modifying the IP address and port number, so the header block have to include the transport layer header. In the transport layer header, my switch should be able to handle both TCP and UDP packet, so I use the header\_union to include both TCP and UDP header, and encapsulate them in the l4\_t header. I also set up three metadata variables to store the length of the TCP, UDP, and total length of the packet, these variables will be used in checksum update, and the variables will be initialized in the parser.

## Parser

```

parser MyParser(packet_in packet,
                out headers hdr,
                inout metadata meta,
                inout standard_metadata_t standard_metadata) {

    state start {
        transition parse_ethernet;
    }

    state parse_ethernet {
        packet.extract(hdr.ethernet);
        transition select(hdr.ethernet.etherType) {
            TYPE_IPV4: parse_ipv4;
            default: accept;
        }
    }

    state parse_ipv4{
        packet.extract(hdr.ipv4);
        transition select(hdr.ipv4.protocol) {
            TYPE_UDP   : parse_udp;
            TYPE_TCP   : parse_tcp;
            default     : accept;
        }
    }

    state parse_udp{
        packet.extract(hdr.l4.udp);
        meta.udp_length = hdr.ipv4.totalLen-20;
        transition accept;
    }

    state parse_tcp{
        packet.extract(hdr.l4.tcp);
        meta.tcp_length = (bit<16>)hdr.l4.tcp.data_offset * 4;
        meta.tot_length = hdr.ipv4.totalLen-20;
        transition accept;
    }

}

```

The parser is a FSM that will parse the packet from the Ethernet header to the transport layer header. The parser will first parse the Ethernet header, then parse the IPv4 header. After parsing the IPv4 header, the parser will select the protocol type, if the protocol is UDP, the parser will parse the UDP header, and store the length of the UDP packet in the metadata. Or the TCP header will be parsed, and the length of the TCP packet and the total length of the packet will be stored in the metadata.

Generally, TCP header includes the option field, but for NAT implementation, it is not necessary to parse the option field. Therefore, I parse TCP field as a fix length field, and treat the option field as payload in the TCP packet.

## Ingress

Ingress part is the main part of the switch, which NAT is implemented in this part. At first, I set up 4 registers to store the necessary information for the NAT translation.

```
157  control MyIngress(inout headers hdr,
158                      inout metadata meta,
159                      inout standard_metadata_t standard_metadata) {
160      register<bit<16>> (max_port_num) port_counter;
161      register<bit<32>> (max_port_num) out_in_IP;
162      register<bit<16>> (max_port_num) out_in_port;
163      register<bit<16>> (max_port_num) ip_port_to_out;
164      action drop() {
165          // drop the packet
166          mark_to_drop(standard_metadata);
167      }
168      action multicast() {
169          standard_metadata.mcast_grp = 1;
170      }
171      action ipv4_forward(bit<48> dstAddr, bit<9> port) {
172          hdr.ethernet.dstAddr = dstAddr;
173          standard_metadata.egress_spec = port;
174      }
175      table ipv4_lookup{
176          key = {
177              hdr.ipv4.dstAddr: lpm;
178          }
179          actions = {
180              ipv4_forward;
181              drop;
182              multicast;
183          }
184          size = 1024;
185          default_action = multicast;
186      }
```

- `port_counter` is used to assign a unique port number for the packet from the internal network to the outer network.
- `out_in_IP` is used to store the mapping between the unique port number and the source IP address.
- `out_in_port` is used to store the mapping between the unique port number and the source port number.
- `ip_port_to_out` is used to store the mapping between the source IP address and the source port number to the unique port number.

For the IP forward part, I set up 3 actions:

- `drop` is used to drop the packet, which set the flag drop in the `standard_metadata`.
- `multicast` is used to set the multicast group in the `standard_metadata`, this action is for use in ARP.
- `ipv4_forward` is used to forward the packet to the specific port, which set the destination MAC address and the egress port in the `standard_metadata`, the table has been set up in `s1-runtime.json`. These three actions are used in the `ipv4_lookup` table, which use longest prefix match to forward the packet to the specific port.

In the apply part, I first using bit operation to check if the IP address is in the internal network.

If the packet's destination IP address is in the internal network, the packet will be forwarded directly to the specific port using the `ipv4_lookup` table.

- This is a simple way to check if the IP address is in the internal network, in a multi-subnet network senario, I think this idea is still useful, but the implementation will be more complex.

```
if((hdr.ipv4.dstAddr>>8)==0x0a0001){
    ipv4_lookup.apply();
}
```

If the destination IP address is not in the internal network, the packet will be checked if it is a TCP or UDP packet, and set up some variables to store the source and destination IP address and port number for the later use in NAT translation.

```

else if(hdr.l4.tcp.isValid() || hdr.l4.udp.isValid()){
    bit<9> src_port=standard_metadata.ingress_port;
    bit<32> index;
    bit<32> base=0;
    bit<32> src_ip=hdr.ipv4.srcAddr;
    bit<32> dst_ip=hdr.ipv4.dstAddr;
    bit<16> src_tcp_port;
    bit<16> dst_tcp_port;
    bit<16> map_in_port;
    if(hdr.l4.tcp.isValid()){
        src_tcp_port=hdr.l4.tcp.srcPort;
        dst_tcp_port=hdr.l4.tcp.dstPort;
    }
    else{
        src_tcp_port=hdr.l4.udp.srcPort;
        dst_tcp_port=hdr.l4.udp.dstPort;
    }
}

```

Later, check if this packet is sent or received from the specific port, if so, the source IP address will be translated to the public IP address, and the port number will not be changed. So the port forwarding part is implemented.

- I think this part can be written in a more concise way, but now I just want to demonstrate this function in a simple way. Sorry~

```

if((src_port<3) && src_tcp_port==80){
    hdr.ipv4.srcAddr=PUB_IP;
}
else if(src_port>2 && dst_tcp_port==80){
    hdr.ipv4.dstAddr=h1_IP;
    hdr.l4.tcp.dstPort=80;
}

```

In the next, the case if the packet is ingress from the internal network and not sent from the specific port. The switch will first put source IP address and port number together into a hash value, then check if the hash value is in the register `ip_port_to_out`. If the corresponding port number is 0, which means this is a new {IP, port} pair, the switch will assign a unique port number to the packet, and store the mapping IP

address and source port number to register `out_in_IP` and `out_in_port`, and store the mapping between the hash value and the unique port number to the register `ip_port_to_out`. Then the source IP address will be translated to the public IP address, and the source port number will be translated to the unique port number. If the corresponding port number is not 0, which means this {IP,port} pair has been assigned a unique port number, the switch will translate the source IP address to the public IP address, and use this unique port number as the source port number. By the way, if the port number is larger than 65500, the port number will be reset to 10000.

```
else if(src_port<3){
    hash(index,HashAlgorithm.crc32,base,{src_ip,src_tcp_port},max_port_num-1);
    bit<16> map_out_port;
    ip_port_to_out.read(map_out_port,index);
    if(map_out_port==0){//new {ip,port} pair
        bit<16> next_port;
        port_counter.read(next_port,0);
        if(next_port==0){
            next_port=10000;
            map_out_port=next_port;
            port_counter.write(0,next_port+1);
            register_modify(index,map_out_port,src_ip,src_tcp_port);
        }
        else{
            map_out_port=next_port;
            if(next_port<65500){
                port_counter.write(0,next_port+1);
            }
            else{
                port_counter.write(0,10000);
            }
            register_modify(index,map_out_port,src_ip,src_tcp_port);
        }
    }
    else{
        hdr.ipv4.srcAddr=PUB_IP;
    }
    if(hdr.l4.tcp.isValid()){
        hdr.l4.tcp.srcPort=map_out_port;
    }
    else{
        hdr.l4.udp.srcPort=map_out_port;
    }
}
```

The last part is the packet is received from the outer network, the switch will check if the destination IP address is the public IP address, if so, the NAT translation will be applied, the switch will use the unique port number to find the corresponding source IP address and port number, and translate the destination port number to the source port number.

```

else if(src_port>2){
    if(hdr.ipv4.dstAddr==PUB_IP){
        out_in_IP.read(hdr.ipv4.dstAddr,(bit<32>)dst_tcp_port);
        out_in_port.read(map_in_port,(bit<32>)dst_tcp_port);
        if(hdr.l4.tcp.isValid()){
            hdr.l4.tcp.dstPort=map_in_port;
        }
        else{
            hdr.l4.udp.dstPort=map_in_port;
        }
    }
}
}
ipv4_lookup.apply();

```

## Checksum update

Checksum update is used after the NAT translation, because the IP address and port number are changed, the checksum in the packet should be updated. In this part, three cases of checksum are needed to be checked and updated:

### 1. IPv4 checksum

IPv4 checksum is calculated by the header of the IPv4, so just simply put the IPv4 header into the checksum update function.

```

update_checksum(
    hdr.ipv4.isValid(),
    { hdr.ipv4.version,
      hdr.ipv4.ihl,
      hdr.ipv4.diffserv,
      hdr.ipv4.totalLen,
      hdr.ipv4.identification,
      hdr.ipv4.flags,
      hdr.ipv4.fragOffset,
      hdr.ipv4.ttl,
      hdr.ipv4.protocol,
      hdr.ipv4.srcAddr,
      hdr.ipv4.dstAddr },
    hdr.ipv4.hdrChecksum,
    HashAlgorithm.csum16);
update_checksum_with_payload(

```

### 2. UDP checksum

UDP checksum is more complex than IPv4 checksum, because the UDP checksum is calculated by the pseudo header, UDP header, and the payload. The pseudo header is the source IP address, destination IP address, protocol type, and the length of the UDP packet. All of them should align in 16 bits, so first in the pseudo header part, 8-bit zero should be combined with the protocol type to align

the 16 bits. Then using the function `update_checksum_with_payload` to update the checksum with the payload we do not parse.

```
update_checksum_with_payload(  
    hdr.l4.udp.isValid(),  
    {  
        //tcp checksum is usually calculated with the following fields  
        //pseudo header+tcp header+tcp payload  
        hdr.ipv4.srcAddr,  
        hdr.ipv4.dstAddr,  
        8w0,           //zero padding with protocol  
        hdr.ipv4.protocol,  
        hdr.l4.udp.length, // 16 bit of tcp length + payload length in bytes  
        hdr.l4.udp.srcPort,  
        hdr.l4.udp.dstPort,  
        hdr.l4.udp.length  
    }, hdr.l4.udp.checksum, HashAlgorithm.csum16);
```

### 3. TCP checksum

TCP checksum is the most difficult part in checksum updating, like UDP checksum, TCP checksum is calculated by the pseudo header, TCP header, and the payload. The pseudo header is the same as the UDP checksum, just pay attention that the length here is the length of the total TCP header with payload. However, in checksum\_update part, any variables cannot be calculated in this part, so I use the metadata calculated in the parser part to update the checksum in the ingress part. Once the metadata is make sure to be correct, the remaining part is the same as the UDP checksum update.

```
hdr.l4.tcp.isValid(),  
{  
    //tcp checksum is usually calculated with the following fields  
    //pseudo header+tcp header+tcp payload  
    hdr.ipv4.srcAddr,  
    hdr.ipv4.dstAddr,  
    8w0,           //zero padding with protocol  
    hdr.ipv4.protocol,  
    meta.tot_length, // 16 bit of tcp length + payload length in bytes  
    hdr.l4.tcp.srcPort,  
    hdr.l4.tcp.dstPort,  
    hdr.l4.tcp.seq_num,  
    hdr.l4.tcp.ack_num,  
    hdr.l4.tcp.data_offset,  
    hdr.l4.tcp.reserved,  
    hdr.l4.tcp.ctl_flag,  
    hdr.l4.tcp.window_size,  
    hdr.l4.tcp.urgent_num  
}, hdr.l4.tcp.checksum, HashAlgorithm.csum16);
```

Conclusion: